

Lab 1: Stampede User Environment

Allocations:

TACC->TRAINING-HPC

XSEDE->TG-TRA140011

May, 2015

Antia Lamas-Linares



Part 0 – Grab the Lab Files

- Login to Stampede

```
$ ssh <username>@stampede.tacc.utexas.edu
```

- Change to your \$WORK directory:

```
$ cdw
```

```
$ pwd
```

```
$ module list
```

- Untar the file lab1_env.tar file (in ~train00) into your directory:

```
$ tar xvf ~train00/stampede_user_env_lab.tar
```

- Move into the newly created lab1_env directory:

```
$ cd stampede_user_env_lab
```

```
$ pwd
```

```
$ ls
```

Part 1 – Run a sanitycheck

- Load the “sanitytool” module
\$ **module load sanitytool**
- Run a sanitycheck in silent mode
\$ **sanitycheck**
- Run a sanitycheck in verbose mode
\$ **sanitycheck -v**

You expect to see some thing like:

```
sanitycheck Version: 1.0.1 @git@ 2014-05-09 20:00
```

```
.....
```

```
All tests passed
```

Part 2 – Using Modules

- List current loaded modules
\$ **module list**
- List available modules
\$ **module av**
- List the help information of one module
\$ **module help hdf5** #or any modules you are interested
- Load extra modules
\$ **module load fftw2 matlab** #or any other modules you need
- Save to your default modules
\$ **module sd**

Use the command “**module reset**” to restore to system defaults.

Part 3 – Run an MPI Batch Job (sbatch)

- Compile the mpipi program:

```
$ mpicc mpipi.c -o mpipi
```

- Open the batch script in an editor to see if you need to change it:

```
$ nano mybatch      # or vi, or emacs
```

<< Change the project “-A” if necessary>>

- Launch the batch job

```
$ sbatch mybatch
```

- Monitor the job’s status (when done, command will return nothing):

```
$ squeue -u <username>
```

```
$ showq | less      # hit space bar to advance
```

```
$ squeue | less     # hit space bar to advance
```

- When job completes, take a look at results:

```
$ ls                # Note presence/names of output files
```

```
$ less mpipi.xxxxx.out # "xxxxx" is your job's jobid
```

```
$ less mpipi.xxxxx.err # "xxxxx" is your job's jobid
```

Part 4– An Interactive Session (idev)

- Compile hello.F90 with and without Open MP.

```
$ ifort -openmp hello.F90 -o hello1
```

```
$ ifort hello.F90 -o hello2
```

- Use idev:

```
$ idev -A TRAINING-HPC
```

- Run the code:

```
$ ./hello1      # you're on a compute node, not a login node
```

```
$ ./hello2
```

- Set OpenMP threads and try again:

```
$ export OMP_NUM_THREADS=4
```

```
$ ./hello1
```

```
$ ./hello2
```

- Clear the thread count and try again:

```
$ unset OMP_NUM_THREADS
```

```
$ ./hello1
```

```
$ ./hello2
```

Part 5 (advanced) – Run MIC App from the Host

- Open a new terminal and recompile to produce "native MIC" code (compilers are not visible from the MIC):

```
$ ifort -mmic -openmp hello.F90 -o hello.mic
```

- Back to the previous terminal, launch the MIC code from the host:

On the Computenode:

```
$ ./hello.mic
```

Note: the program reports 244 “processors” because each MIC core has four hardware threads. It may not be efficient to run this many threads.

- From the host, modify the MIC thread count and try again:

```
$ export MIC_ENV_PREFIX=MIC
```

```
$ ./hello.mic
```

Part 6 (advanced) – Visit the MIC

- First note the full path to your working directory:
`$ echo $WORK # you'll need this info when you get to the MIC`
- Go the MIC using ssh:
`$ ssh mic0 # the "zero" identifies the MIC card`
- Move into the stampede_user_env directory with explicit cd:
`$ cd /work/12345/yourname #replace with your own path`
`$ cd stampede_user_env_lab`
- Run your MIC code:
`$./hello.mic`
- It fails because it cannot find libiomp5.so.
`$ newpath="/opt/apps/intel/13/composer_xe_2013.2.146/compiler/lib/mic"`
`$ export LD_LIBRARY_PATH=$newpath:${LD_LIBRARY_PATH}`
`$./hello.mic`
- End session:
`$ exit # to return to host`
`$ exit # to end idev session`