

MPI lecture and labs 5

Victor Eijkhout

2016

Sub-computations

Sub-computations

Simultaneous groups of processes, doing different tasks, but loosely interacting:

- Simulation pipeline: produce input data, run simulation, post-process.
- Climate model: separate groups for air, ocean, land, ice.
- Quicksort: split data in two, run quicksort independently on the halves.
- Processor grid: do broadcast in each column.

New communicators are formed recursively from `MPI_COMM_WORLD`.

Communicator duplication

Simplest new communicator: identical to a previous one.

```
int MPI_Comm_dup(MPI_Comm comm, MPI_Comm *newcomm)
```

This is useful for library writers:

```
MPI_Isend(...); MPI_Irecv(...);  
// library call  
MPI_Waitall(...);
```

Use of a library

```
library my_library(comm);  
MPI_Isend(&sdata,1,MPI_INT,other,1,comm,&(request[0]));  
my_library.communication_start();  
MPI_Irecv(&rdata,1,MPI_INT,other,MPI_ANY_TAG,  
          comm,&(request[1]));  
MPI_Waitall(2,request,status);  
my_library.communication_end();
```

Use of a library

```
int library::communication_start() {  
    int sdata=6,rdata;  
    MPI_Isend(&sdata,1,MPI_INT,other,2,comm,&(request[0]));  
    MPI_Irecv(&rdata,1,MPI_INT,other,MPI_ANY_TAG,  
             comm,&(request[1]));  
    return 0;  
}  
  
int library::communication_end() {  
    MPI_Status status[2];  
    MPI_Waitall(2,request,status);  
    return 0;  
}
```

Wrong way

```
// commdup_wrong.cxx
class library {
private:
    MPI_Comm comm;
    int mytid, ntids, other;
    MPI_Request *request;
public:
    library(MPI_Comm incomm) {
        comm = incomm;
        MPI_Comm_rank(comm, &mytid);
        other = 1-mytid;
        request = new MPI_Request[2];
    };
    int communication_start();
    int communication_end();
};
```

Right way

```
// commdup_right.cxx
class library {
private:
    MPI_Comm comm;
    int mytid, ntids, other;
    MPI_Request *request;
public:
    library(MPI_Comm incomm) {
        MPI_Comm_dup(incomm, &comm);
        MPI_Comm_rank(comm, &mytid);
        other = 1-mytid;
        request = new MPI_Request[2];
    };
    ~library() {
        MPI_Comm_free(&comm);
    }

    int communication_start();
};
```


Disjoint splitting

Split a communicator in multiple disjoint others.

Give each process a 'colour', group processes by colour:

```
int MPI_Comm_split(MPI_Comm comm, int color, int key,  
                  MPI_Comm *newcomm)
```

Row/column example

```
MPI_Comm_rank( MPI_COMM_WORLD, &mytid );  
proc_i = mytid % proc_column_length;  
proc_j = mytid / proc_column_length;  
  
MPI_Comm column_comm;  
MPI_Comm_split( MPI_COMM_WORLD, proc_j, mytid, &column_comm );  
  
MPI_Bcast( data, ... column_comm );
```


Exercise 1

Organize your processors in a grid, and make subcommunicators for the rows and columns. Do a broadcast from the first row and column through the columns and rows respectively.

If you let the broadcast value be the column/row number, then processor (i,j) winds up with the numbers i and j . Test this.

Exercise 2

Implement a recursive algorithm for matrix transposition:



- Swap blocks $(1,2)$ and $(2,1)$; then
- Divide the processors into four subcommunicators, and apply this algorithm recursively on each;
- If the communicator has only one process, transpose the matrix in place.


- Non-disjoint subcommunicators through process groups.
- Intra-communicators and inter-communicators.
- Process topologies: cartesian and graph.

One-sided communication

With two-sided messaging, you can not just put data on a different processor: the other has to expect it and receive it.

- Sparse matrix: it is easy to know what you are receiving, not what you need to send. Usually solved with complicated preprocessing step.
- Neuron simulation: spiking neuron propagates information to neighbours. Uncertain when this happens.
- Other irregular data structures: linked lists, hash tables.

One-sided concepts



- A process has a *window* that other processes can access.
- Origin: process doing a one-sided call; target: process being accessed.
- One-sided calls: `MPI_Put`, `MPI_Get`, `MPI_Accumulate`.
- Various synchronization mechanisms.


Active target synchronization

All processes call `MPI_Win_fence`. Epoch is between fences:

```
MPI_Win_fence(MPI_MODE_NOPRECEDE, win);  
if (mytid==producer)  
    MPI_Put( /* operands */, win);  
MPI_Win_fence(MPI_MODE_NOSUCCEED, win);
```

Second fence indicates that one-sided communication is concluded: target knows that data has been put.

Window creation



```
MPI_Win_create (void *base, MPI_Aint size,  
               int disp_unit, MPI_Info info,  
               MPI_Comm comm, MPI_Win *win)
```

- size: in bytes
- disp_unit: sizeof(type)
- Also: MPI_Win_allocate, can use dedicated fast memory.

C:

```
int MPI_Put(  
    const void *origin_addr, int origin_count, MPI_Datatype origin_datatype,  
    int target_rank, MPI_Aint target_disp, int target_count, MPI_Datatype target_datatype,  
    MPI_Win win)
```

Semantics:

IN origin_addr: initial address of origin buffer (choice)
IN origin_count: number of entries in origin buffer (non-negative integer)
IN origin_datatype: datatype of each entry in origin buffer (handle)
IN target_rank: rank of target (non-negative integer)
IN target_disp: displacement from start of window to target buffer (non-negative integer)
IN target_count: number of entries in target buffer (non-negative integer)
IN target_datatype: datatype of each entry in target buffer (handle)
IN win: window object used for communication (handle)

Fortran:

```
MPI_Put(origin_addr, origin_count, origin_datatype,  
    target_rank, target_disp, target_count, target_datatype, win, ierror)  
TYPE(*), DIMENSION(..), INTENT(IN), ASYNCHRONOUS :: origin_addr  
INTEGER, INTENT(IN) :: origin_count, target_rank, target_count  
TYPE(MPI_Datatype), INTENT(IN) :: origin_datatype, target_datatype
```

Exercise 3

Write code where process 0 randomly writes in the window on 1 or 2.


```
// randomput_sk1.c
MPI_Win_create(&window_data, sizeof(int), sizeof(int),
               MPI_INFO_NULL, comm, &the_window);

for (int c=0; c<10; c++) {
    float randomfraction = (rand() / (double)RAND_MAX);
    if (randomfraction>.5)
other = 2;
    else other = 1;
    window_data = 0;
    your_code_goes_here.....
    my_sum += window_data;
}

if (mytid>0 && mytid<3)
    MPI_Send(&my_sum, 1, MPI_INT, other, 0, comm);
    MPI_Configure(&my_sum, 1, MPI_INT, mytid, my_sum);
}
```

A second active synchronization

Use `Post`, `Wait`, `Start`, `Complete` calls



More fine-grained than fences.

Passive target synchronization

Lock a window on the target:

```
MPI_Win_lock (int locktype, int rank, int assert, MPI_Win win)
MPI_Win_unlock (int rank, MPI_Win win)
```

Atomic operations:

```
int MPI_Fetch_and_op(const void *origin_addr, void *result_addr,
    MPI_Datatype datatype, int target_rank, MPI_Aint target,
    MPI_Op op, MPI_Win win)
```

```

// passive.cxx
if (mytid==repository) {
    // Processor zero creates a table of inputs
    // and associates that with the window
}
if (mytid!=repository) {
    float contribution=(float)mytid,table_element;
    int loc=0;
    MPI_Win_lock(MPI_LOCK_EXCLUSIVE,repository,0,the_window);
    // read the table element by getting the result from adding
    err = MPI_Fetch_and_op(&contribution,&table_element,MPI_FLOAT,
                          repository,loc,MPI_SUM,the_window); CHK(err);
    MPI_Win_unlock(repository,the_window);
}

```