

REVIEW: Worksharing Construct - DO/for Directive

Schedule Clause In a Nutshell

- `schedule(static)` n iterations divided in blocks of n/p .
- `schedule(static,m)` iterations divided in blocks of m ('chunk size'), assigned cyclically.
- `schedule(dynamic)` single iterations, assigned whenever a thread is idle
- `schedule(dynamic,m)` blocks of m iterations, assigned whenever a thread is idle
- `schedule(guided)` decreasing size blocks
- `schedule(auto)` leave it up to compiler/runtime
- `schedule(runtime)` using environment variable `OMP_SCHEDULE`

Worksharing Construct - Reduction a brief (very brief) intro

```
int array[8] = { 1, 1, 1, 1, 1, 1, 1, 1};  
int sum = 0, i;
```

```
#pragma omp parallel for reduction(+:sum)  
for (i = 0; i < 8; i++)  
{  
    sum += array[i];  
}
```

```
printf("total %d\n", sum);
```

- Reductions are atomic operations
- Can be solved by private variable per thread
- reduction clause is shorthand for all that

An operation acting on shared memory is **atomic** if it completes in a single step relative to other threads. When an atomic store is performed on a shared variable, no other thread can observe the modification half-complete. When an atomic load is performed on a shared variable, it reads the entire value as it appeared at a single moment in time.

Exercise 3 - Code

Take 20 minutes and enhance this code with OpenMP.

```
#include <stdlib.h>
#include <stdio.h>
#include <math.h>
#include <omp.h>

int main(int argc, char **arg) {

    int nsteps=1000000000; // that's one
                           billion

    double tstart, tend, elapsed,
           pi, quarterpi, h;

    int i;

    tstart = omp_get_wtime(); //gettime();

    quarterpi = 0.; h = 1./nsteps;
```

```
    for (i=0; i<nsteps; i++)
    {
        double x = i*h, y = sqrt(1-x*x);
        quarterpi += h*y;
    }

    pi = 4*quarterpi;
    tend = omp_get_wtime(); //gettime();
    elapsed = tend-tstart;

    printf("Computed pi=%e in %6.3f seconds\n", pi,
    elapsed);

    return 0;
}
```

Exercise 3 - Enhancements?

- Use the `omp parallel for` construct to parallelize the loop.
- Are you seeing any odd behavior? Or weird discrepancies? See if reduction clause can fix this.
 - More on the reduction clause soon!
- Your code should now see a decent speedup, using up to 8 cores. However, it is possible to get completely linear speedup. For this you need to adjust the schedule.

Exercise 3a - Make it parallel, and a schedule

- add 'adaptive integration': where needed, the program refines the step size. This means that the iterations no longer take a predictable amount of time.

```
for (i=0; i<nsteps; i++) {  
    double x = i*h, x2 = (i+1)*h,  
  
    y = sqrt(1-x*x), y2 = sqrt(1-x2*x2),  
  
    slope = (y-y2)/h;  
  
    if (slope>15) slope = 15;  
  
    int samples = 1+(int)slope,  
  
    is;
```

```
    for (is=0; is<samples; is++)  
    {  
        double hs = h/samples,  
  
        xs = x + is*hs,  
  
        ys = sqrt(1-xs*xs);  
  
        quarterpi += hs*ys;  
  
        nsamples++;  
    }  
}  
  
pi = 4*quarterpi;
```

Exercise 3a - Make it parallel, and add the schedule clause

- Use the `omp parallel for` construct to parallelize the loop. As in the previous lab, you may at first see an incorrect result. Use the reduction clause to fix this.
- Your code should now see a decent speedup, using up to 8 cores. However, it is possible to get completely linear speedup. For this you need to adjust the schedule.
- Start by using `schedule(static,n)`. Experiment with values for `n`. When can you get a better speedup? Explain this.
- Since this code is somewhat dynamic, try `schedule(dynamic)`. This will actually give a fairly bad result. Why? Use `schedule(dynamic,n)` instead, and experiment with values for `n`.
- Finally, use `schedule(guided)`, where OpenMP uses a heuristic. What results does that give?

Ordered Iterations Clause

Let's recall our OpenMP HelloWorld program.

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>

int main (int argc, char *argv[])
{
    int nthreads, tid;

    /* Fork a team of threads giving them their own copies of
       variables */
    #pragma omp parallel private(nthreads, tid)
    {
```

```
        /* Obtain thread number */
        tid = omp_get_thread_num();

        printf("Hello World from thread = %d\n",
               tid);
        /* Only master thread does this */
        if (tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }

    } /* All threads join master thread and disband */
}
```

Did all the threads execute their code at the same time?

Ordered Iterations Clause

What if we added a loop? Would all the loops execute in their proper order?

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>

int main (int argc, char *argv[])
{
    int nthreads, tid;

    /* Fork a team of threads giving them their own copies of
       variables */
    #pragma omp parallel private(nthreads, tid)
    {
        /* Obtain thread number */
        tid = omp_get_thread_num();
```

```
        #pragma omp parallel for
        for (int i=0; i<5)
            printf("Hello World %d from thread = %d\n",
                   i, tid);
    /* Only master thread does this */
    if (tid == 0)
    {
        nthreads = omp_get_num_threads();
        printf("Number of threads = %d\n", nthreads);
    }

    } /* All threads join master thread and disband */
}
```

And they do not... Iterations in a parallel loop that are ran in parallel do not execute in lockstep.

Ordered Iterations Clause

The ordered clause coupled with the ordered directive can force execution in the right order:

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>

int main (int argc, char *argv[])
{
    int nthreads, tid;

    /* Fork a team of threads giving them their own copies of
       variables */
    #pragma omp parallel private(nthreads, tid)
    {
        /* Obtain thread number */
        tid = omp_get_thread_num();
```

```
        #pragma omp parallel for ordered
        {
            for (int i=0; i<5)
                printf("Hello World %d from thread = %d\n",
                    i, tid);
        }
        #pragma omp ordered
        /* Only master thread does this */
        if (tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }

    } /* All threads join master thread and disband */
}
```

NOTE: Each iteration can only encounter *one* ordered directive

Nowait Clause

Recall that during a for clause, threads wait until all threads are finished before continuing. The `nowait` clause allows code to continue with the next line of code in a parallel region.

```
...
    #pragma omp parallel private(nthreads, tid)
    {
/* Obtain thread number */
        tid = omp_get_thread_num();
        #pragma omp parallel for nowait
        {
            for (int i=0; i<1000)
                printf("Hello World %d from thread = %d\n",
                    i, tid);
        }
    }
```

```
#pragma omp parallel for
{
    for (int i=0; i<1000)
        printf("Goodbye Cruel World %d from thread =
            %d\n", i, tid);
    }
}
```

Any thread that completes its task in the first for loop, may now continue to the second for loop. NOTE: This requires both loops to have the same schedule.

Section Construct

- When doing large blocks of independent computations, split the work into sections.
 - Parallel Loops are an example of independent work units which are numbered
 - If you have a set number of work units, use the section construct
 - Section constructs may contain any number of nested section constructs
 - They need to be coded in such a way that any thread on the current team of threads can execute any section or multitude of sections

```
#pragma omp section
{
    #pragma omp section
    {
        // one calculation
    }
    #pragma omp section
    {
        // another calculation
    }
}
```

Sections

An example. Suppose we have $y = f(x) + g(x)$

```
...  
    double y1,y2;  
    #pragma omp sections  
    {  
        #pragma omp section  
        y1 = f(x);  
        #pragma omp section  
        y2 = g(x);  
    }  
    y = y1+y2;
```

Given the above code snippet, how can we add a reduction clause for better efficiency?

Single Directive

Specifies that only a single thread should execute this section of the program.

```
#pragma omp parallel num_threads(2)
{
    #pragma omp single
    // Only a single thread can read the input.
    printf_s("read input\n");

    // Multiple threads in the team compute the
    // results.
    printf_s("compute results\n");

    #pragma omp single
    // Only a single thread can write the output.
    printf_s("write output\n");
}
```

- May or may not be the master thread
- Synchronizes through the implicit barrier
 - it's a **work sharing construct**, there is an implicit barrier after it, which guarantees that all threads have the correct value in their local memory

Master Directive

Specifies that only the master thread should execute a section of the program.

```
#pragma omp parallel
{
    // Perform some computation.
    #pragma omp for
    for (i = 0; i < N; i++)
        a[i] = i * i;

    // Print intermediate results.
    #pragma omp master
    for (i = 0; i < N; i++)
        printf_s("a[%d] = %d\n", i, a[i]);

    // Wait.
    #pragma omp barrier

    // Continue with the computation.
    #pragma omp for
    for (i = 0; i < N; i++)
        a[i] += i;
}
```

- The master directive supports no OpenMP clauses.
- Does not synchronize through the implicit barrier

Exercise 4 - Homework, Matrix Multiplication Revisted

Write a program that performs matrix multiplication.

- Create two double dimensioned arrays, populate them with random numbers using a single thread
- Create a set of nested loops that multiplies the two arrays (your matrices) together in parallel and using a simple reduction clause (we will be covering reduction in depth)
- Add a worksharing clause, you may need to experiment with this.
- add a time function, and record start time and end time of your loop and how long the loop took to process.
- run this for a 10x10, 100x100, and 1000x1000

How are your running times now compared to when we first did this assignment (Exercise 2)?